

Time Series Forecasting with Tree-Based Ensemble Methods: Part II

April 15, 2022

About the Presenter

- Brett Efaw
- Data Scientist at Idaho Power Company
- BEfaw@idahopower.com or brettefaw@gmail.com
- Work with various business groups to implement Data Science solutions
 - Decision support
 - Predictive modeling
 - Forecasting
 - BI and DS reporting/analysis applications



Overview

- Part I (DS Seminar 2022-02-25):
 - [data_science_seminar/2022_1 at master · lselector/data_science_seminar · GitHub](#)
 - Time series forecasting approaches
 - Compare XGBoost to Auto-ARIMA on a simple simulated data set
 - Part II (DS Seminar 2022-04-15):
 - Use a more complex simulated data set
 - Real world data set (from Kaggle)
 - Auto-ARIMA with exogenous regressors
 - Simulated time series data set
 - Multiple cyclic patterns/seasonality
 - Non-linear trend
 - Modeling/forecasting process
 - Extract the “trend” from the time series
 - Model 1: Fit and forecast the trend component
 - Model 2: Use XGBoost (or similar tree-based algo) to model remainder
 - Forecast = combination of model 1 + model 2 predictions
-

(A More Interesting) Simulated Data Set

- Cyclic pattern features:
 - Day of week
 - Day of month
 - Day of year
- Sequence features:
 - Day number
 - Year
- Holiday effects:
 - Holiday
 - Day before/after holiday
- More realistic trend:
 - Non-linear
 - Randomly-spaced up/down steps
 - `numpy.random.choice()`
 - Smoothed out across time
 - `savgol_filter()`


```
df["dw"] = df["dt"].dt.dayofweek+1
df["dm"] = df["dt"].dt.day
df["dy"] = df["dt"].dt.dayofyear
df["dn"] = df.index+1
df["yr"] = df["dt"].dt.year
```

```
cal = calendar()
holidays = cal.holidays(start = df["dt"].min(), end = df["dt"].max())
df["hd"] = df["dt"].isin(holidays).astype(int) # Federal Holiday
df["bh"] = df["hd"].shift(periods = -1, fill_value = 0) # Day Before Holiday
df["ah"] = df["hd"].shift(periods = +1, fill_value = 0) # Day After Holiday
```

```
# Noise for each date
df["noise"] = np.random.normal(0, 1.0, len(df))
# Step changes randomly dispersed throughout the time series
chngepts = sorted(np.random.choice(df["dn"], size=50, replace=False))
df["change"] = np.where(df["dn"].isin(chngepts), 1, 0)
df["step"] = np.random.choice([-1,1], size=len(df), p=[0.5,0.5]) * df["change"]
df["step"] = df["step"].cumsum()
# Smooth step changes (i.e. non-linear trend)
df["trend"] = savgol_filter(df["step"].values, window_length=100, polyorder=1)
```

	dt	dw	dm	dy	dn	yr	sdw	cdw	sdm	cdm	sdyl	cdyl	hd	bh	ah	noise	change	step	trend	y
0	2015-01-01	4	1	1	1	2015	0.433884	-0.900969	0.000000	1.000000	0.000000	1.000000	1	0	0	-1.626860	0	0	-0.217624	101.688431
1	2015-01-02	5	2	2	2	2015	-0.433884	-0.900969	0.201299	0.979530	0.017166	0.999853	0	0	1	0.613607	0	0	-0.210399	99.266203
2	2015-01-03	6	3	3	3	2015	-0.974928	-0.222521	0.394356	0.918958	0.034328	0.999411	0	0	0	-1.501417	0	0	-0.203174	99.445012
3	2015-01-04	7	4	4	4	2015	-0.781831	0.623490	0.571268	0.820763	0.051479	0.998674	0	0	0	-0.344402	0	0	-0.195950	101.743491
4	2015-01-05	1	5	5	5	2015	0.000000	1.000000	0.724793	0.688967	0.068615	0.997643	0	0	0	2.279467	0	0	-0.188725	105.570760

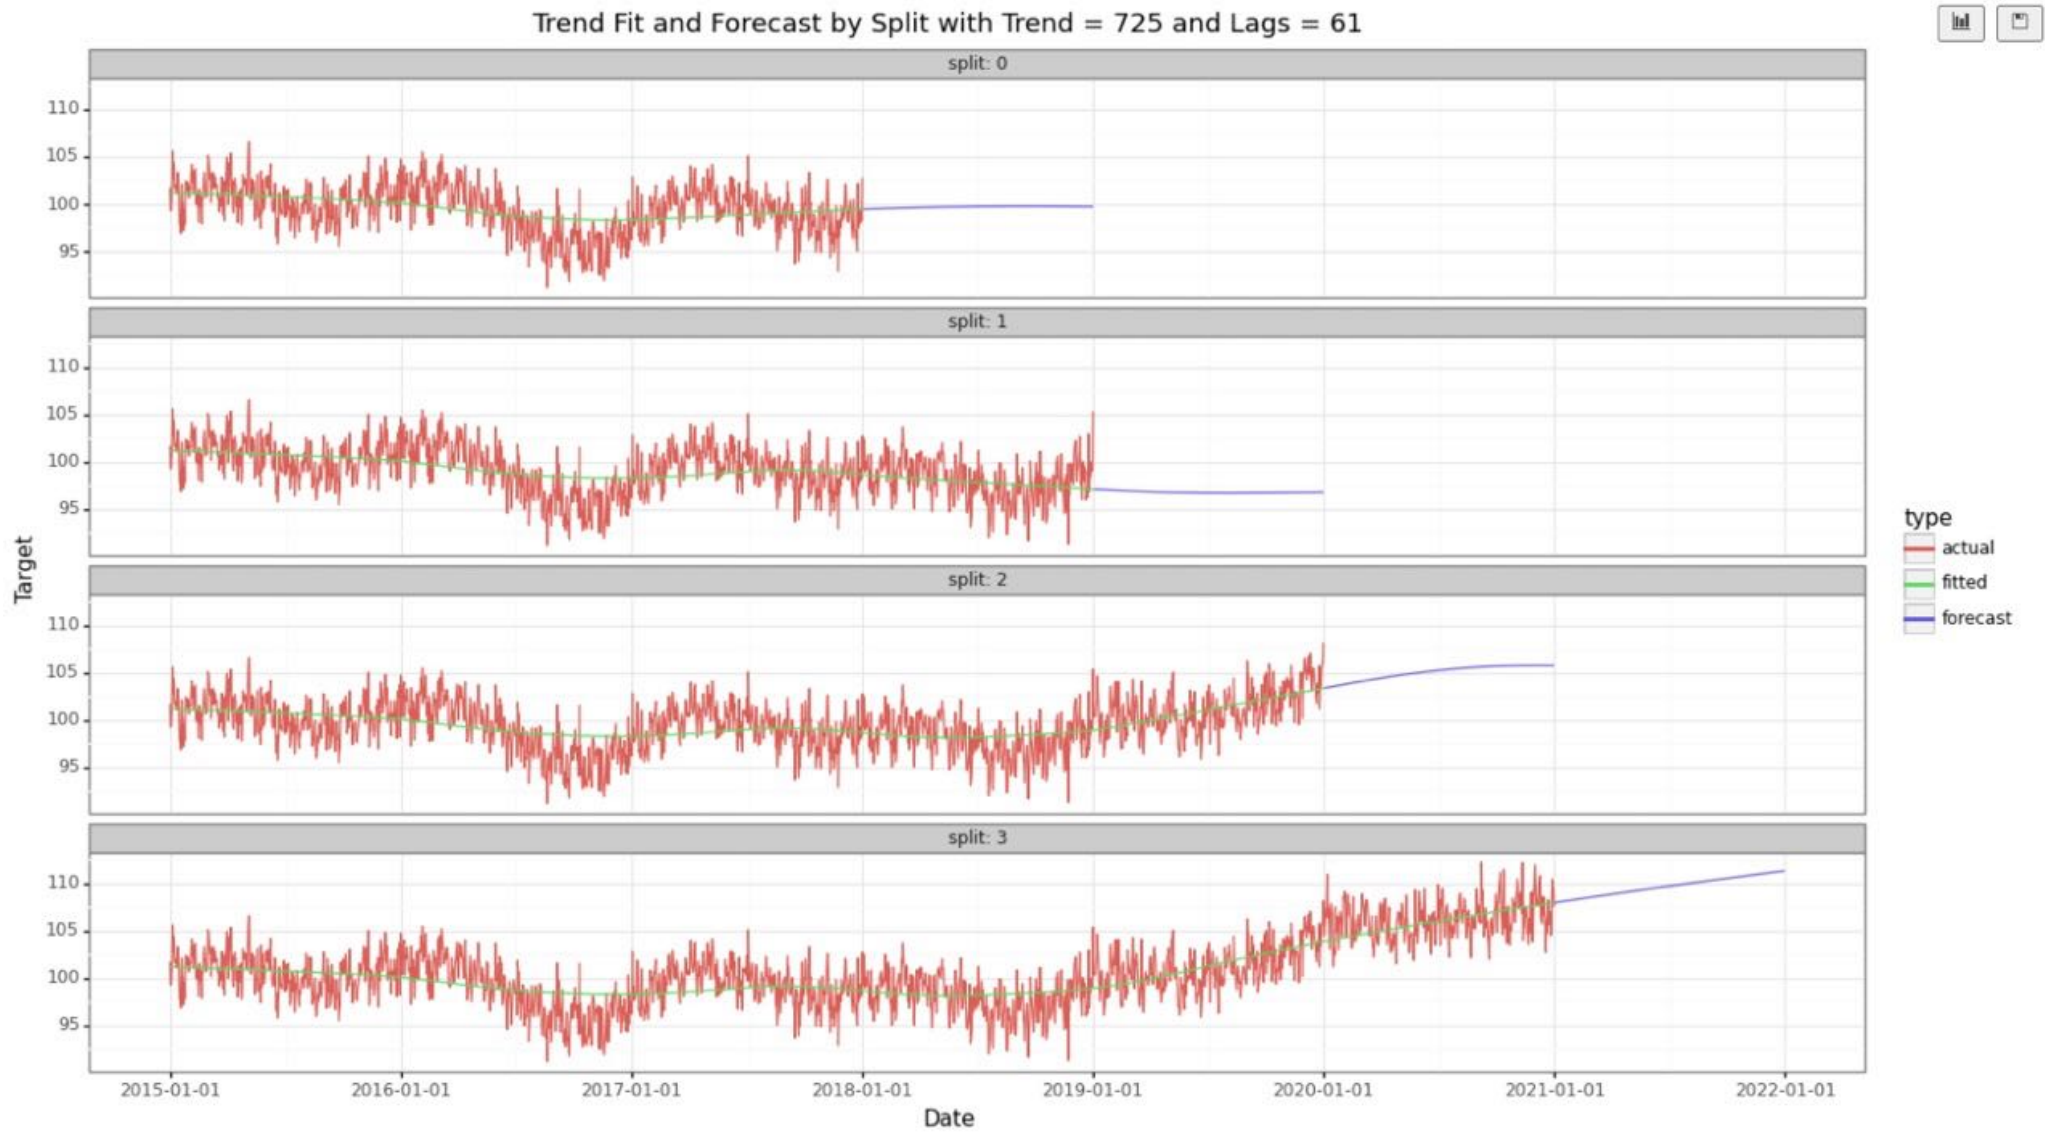
Target Variable




Model 1: Fit and Forecast Trend

- M1-a. Extract the trend
 - Use `statsmodels.tsa.seasonal.STL()` to decompose the time series (i.e. extract the trend)
 - Trend = the number of time steps to include in the decomposition
 - larger values = smoother trend (favors long-term)
 - smaller values = wigglier trend (favors short-term)
 - M1-b. Model the trend
 - Use `statsmodels.tsa.ar_model.AutoReg()` to model (and forecast) the trend
 - Lags = the number of time steps prior to consider as lagged inputs in regression model
 - M1-c. Detrend the past (need these for Model 2)
 - M1-d. Forecast the trend using model in M1-b
 - Other parameters
 - Period = the frequency of the time series (daily = 7, monthly = 12, quarterly = 4, etc.)
 - Number of splits = the number of train/test splits for back-casting / model evaluation
 - Test size = the number of time steps in the forecast horizon
-

Trend Fit and Forecast by Split with Trend = 725 and Lags = 61




De-trended by Split



Model 2: Tree-Based Ensemble

- Use tree-based ensemble to model de-trended time series
 - De-trending ensures the tree-based model can safely predict within the range of training examples.
 - Features are calendar-based:
 - Day-of-week
 - Day-of-month
 - Day-of-year
 - Holidays (adjacent)
 - Month-of-year
 - Etc.
 - HistGradientBoostingRegressor() parameters
 - Max iterations = 300
 - Max tree depth = 30
 - Learning rate = 0.1
 - L2 regularization = 0.1
 - Forecast = trend forecast + tree-based prediction
-

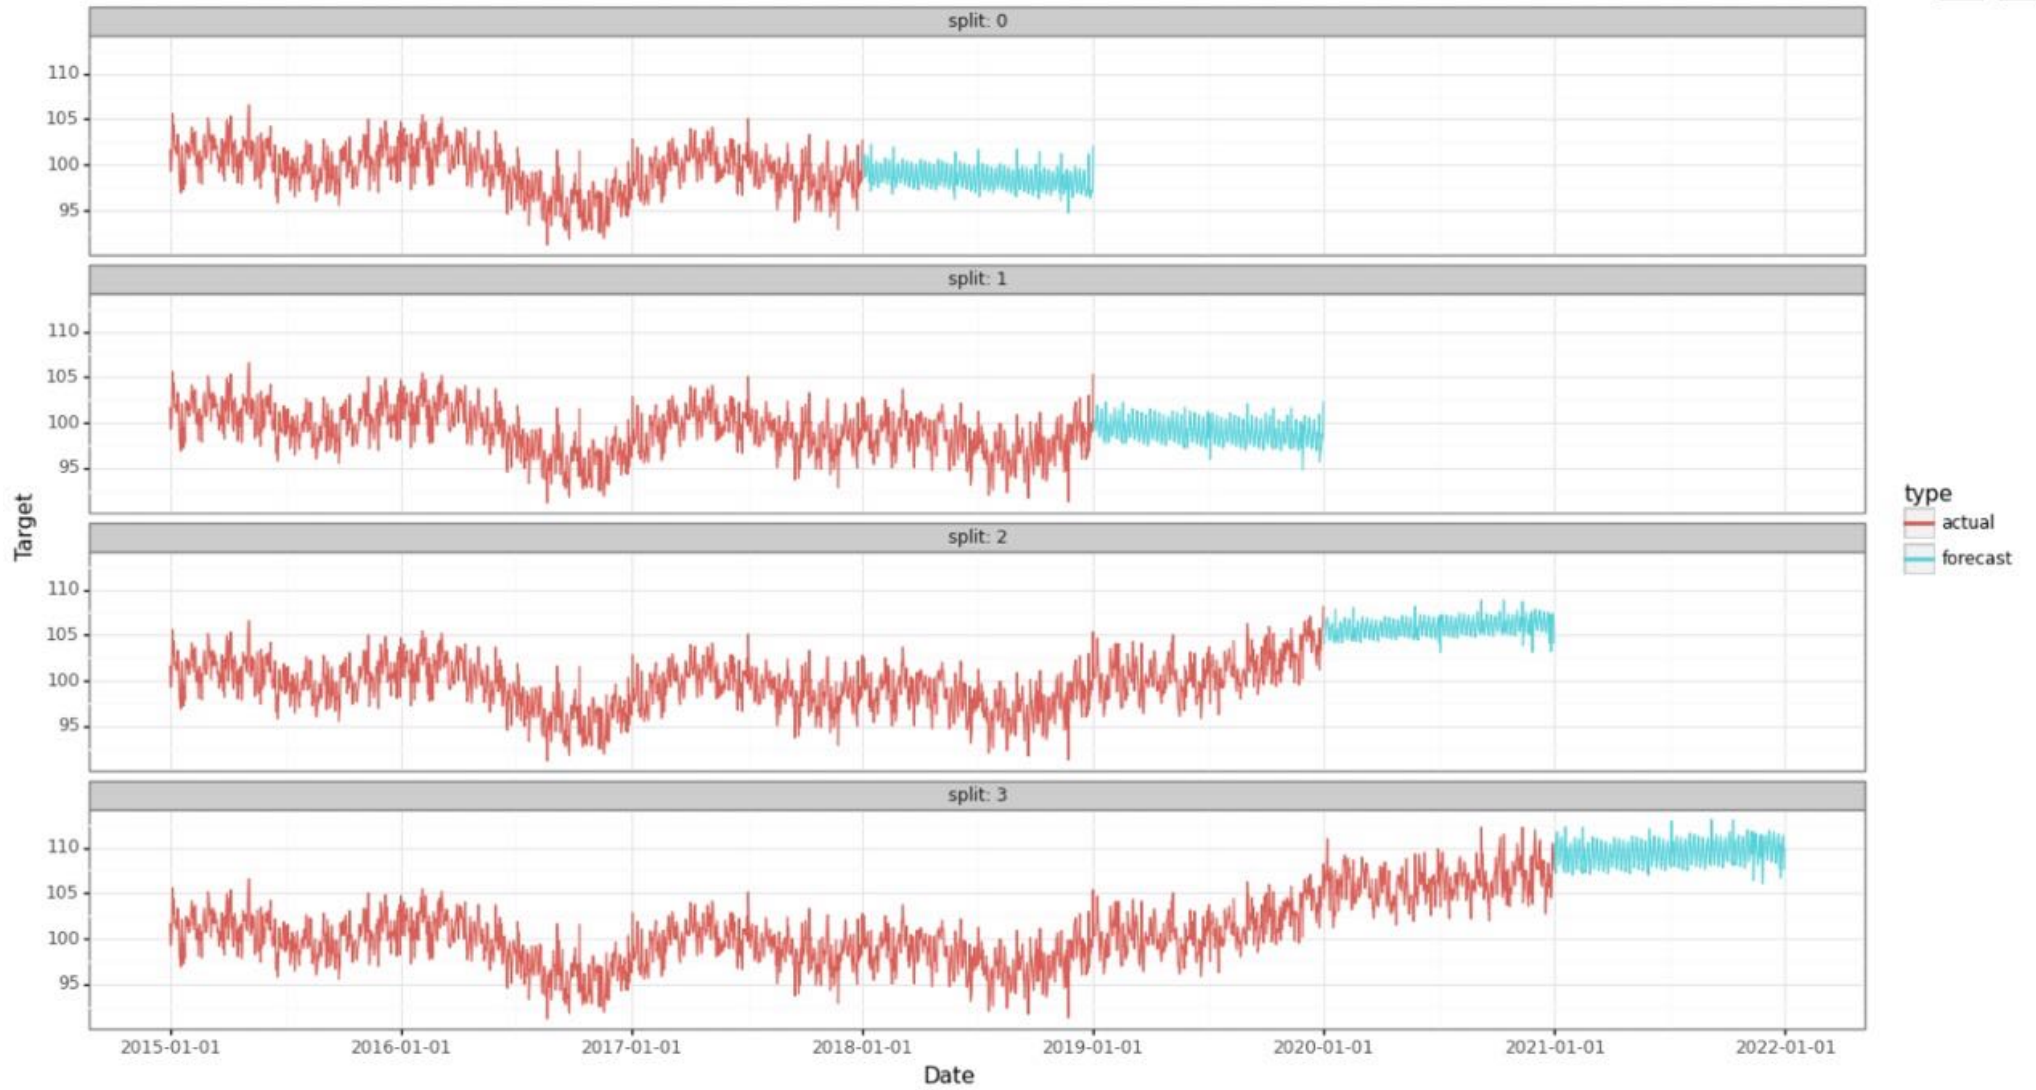
Tree-Model + Trend Forecast by Split




Auto-ARIMA Procedure

- Compare trend + tree-based model to auto-ARIMA procedure
 - Implemented in Python via the pmdarima library
 - based on auto.arima() function from R forecast package
 - SARIMAX = **S**easonal **A**uto-**R**egressive Integrated **M**oving-Average with e**X**ogenous
 - ARIMA, plus:
 - Enable seasonality parameter
 - Include exogenous regressors (calendar features)
 - Parameters
 - D = 1 (difference at lag-1)
 - Period = 7 (daily frequency)
 - Seasonal = True
 - Stepwise = True
 - Exogenous = matrix of external regressors (calendar features)
 - Max iterations = 10 (default = 50)
 - Note: For real-world data sets, try adding Fourier transforms instead of calendar features as exogenous regressors
-

Auto-ARIMA Forecasts by Split



Forecast Comparison



Comparison Between Models on Simulated Data Set


- STL/AutoReg Trend + XGBoost
 - Train Time*:
 - Trend: from 0:00:0.49 to 0:00:0.87
 - De-Trending: 0:00:0.01
 - XGBoost: from 0:00:1.26 to 0:00:1.36
 - Total: ~2-3 seconds per split
 - Accuracy:
 - from 1.3% to 3.2%
 - Avg across splits: 2.3%
- Auto-ARIMA
 - Train Time*:
 - from 1min 16sec to 3min 28sec
 - Accuracy:
 - from 1.4% to 3.1%
 - Avg across splits: 2.1%

*4 splits with 365 test size each

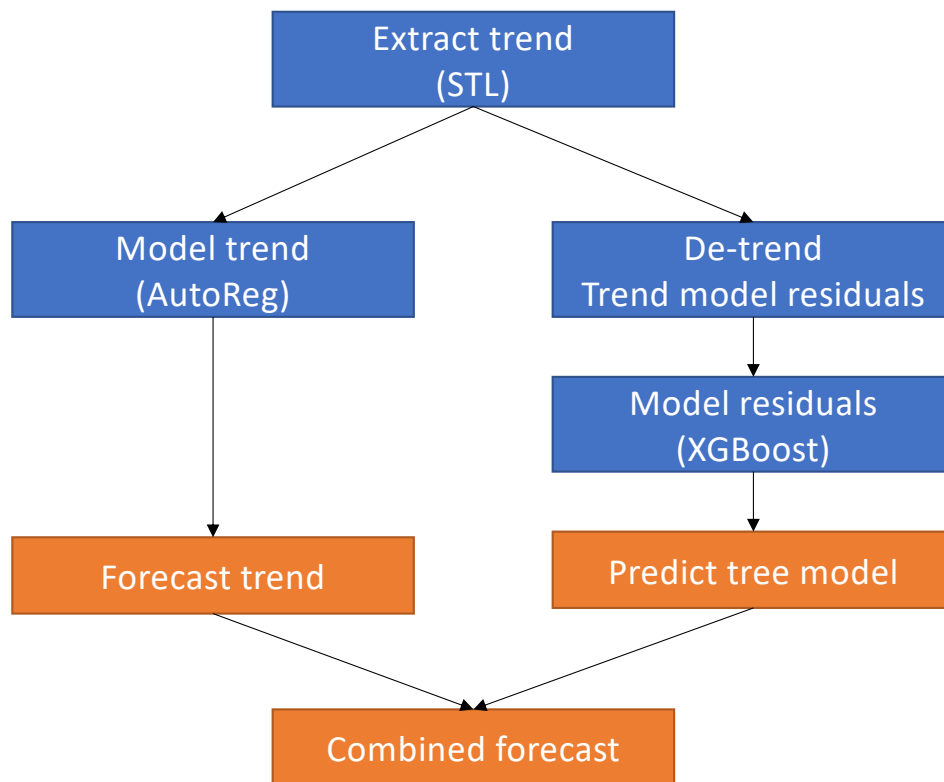
Real-World Data Set

Data Set on Kaggle: [Daily website visitors \(time series regression\) | Kaggle](#)

This file contains 5 years of daily time series data for several measures of traffic on a statistical forecasting teaching notes website whose alias is [statforecasting.com](#). The variables have complex seasonality that is keyed to the day of the week and to the academic calendar. The patterns you see here are similar in principle to what you would see in other daily data with day-of-week and time-of-year effects.



Trend + Tree-Based Modeling Procedure

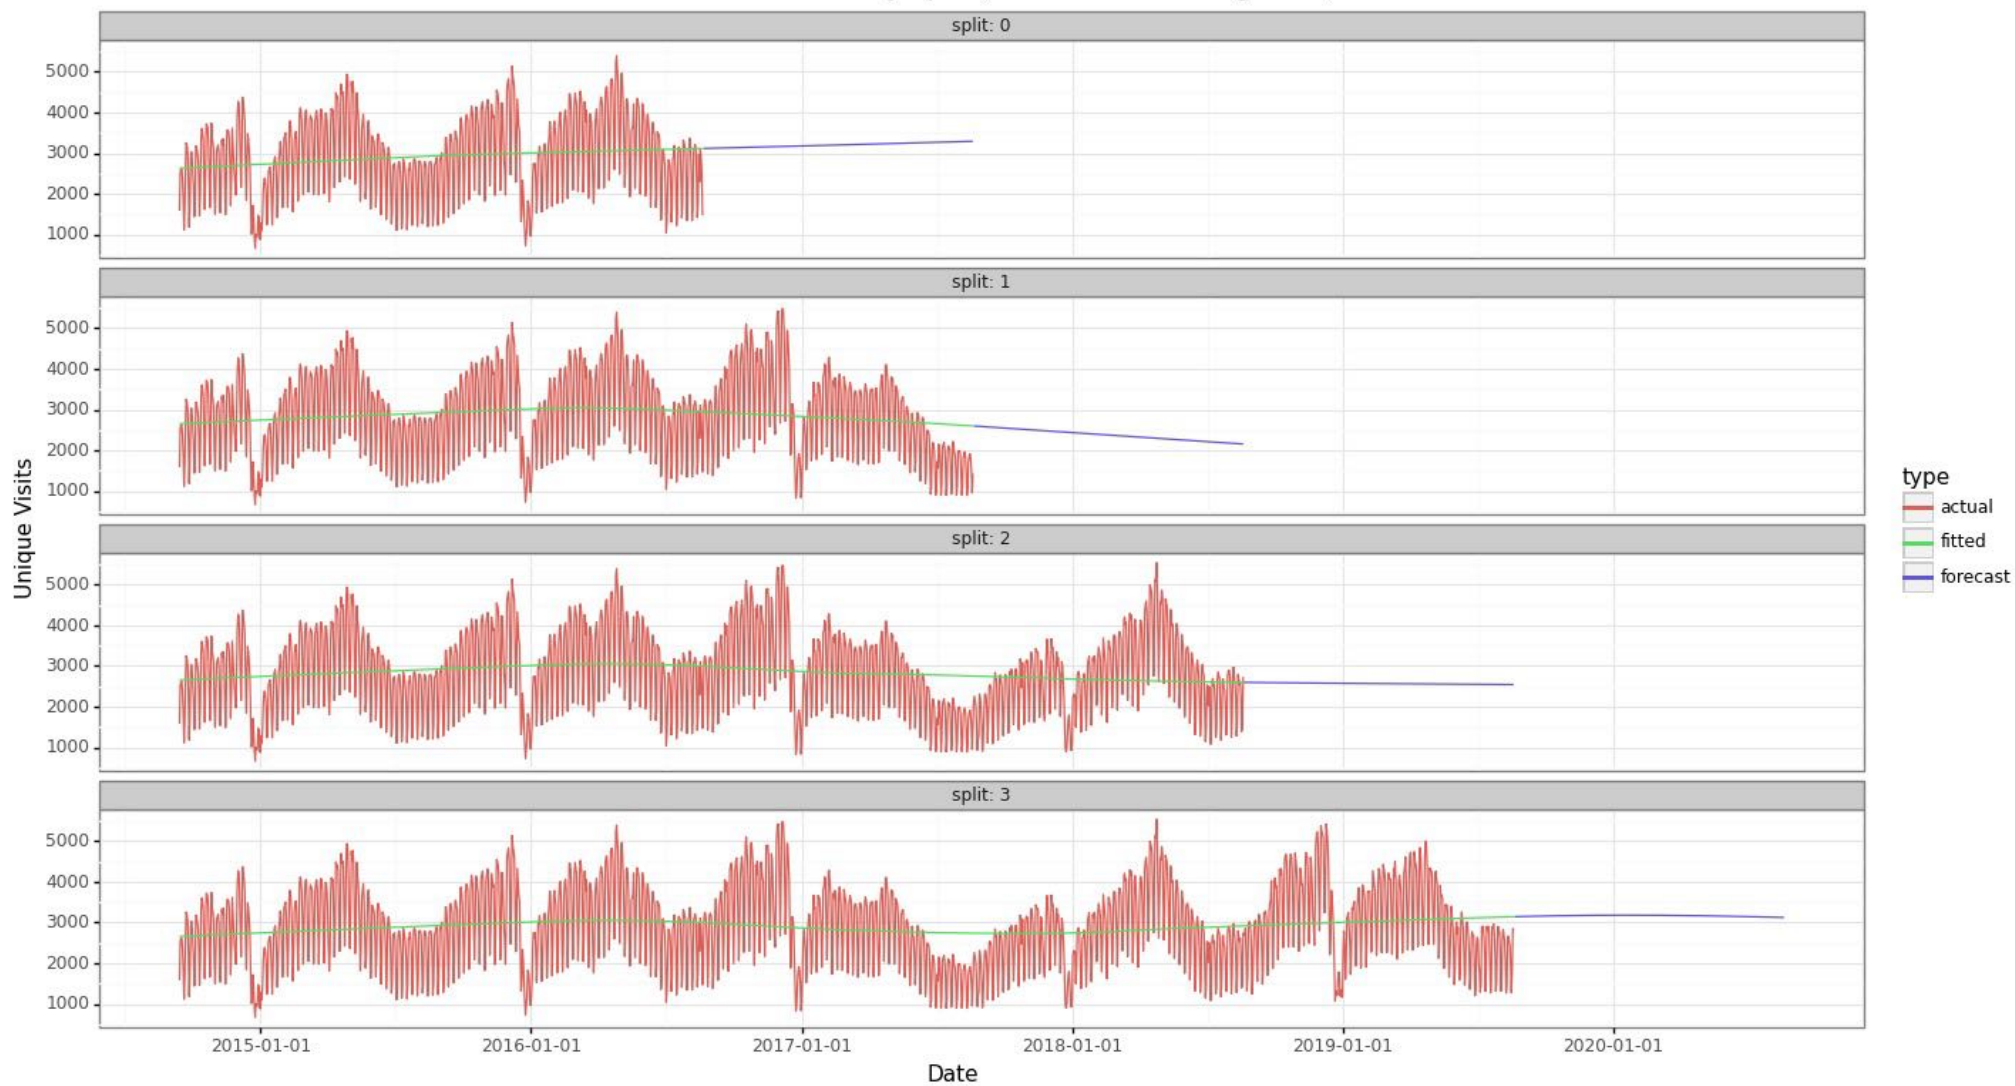


Color Legend


Training Split i.e. Historical

Testing Split i.e. Future

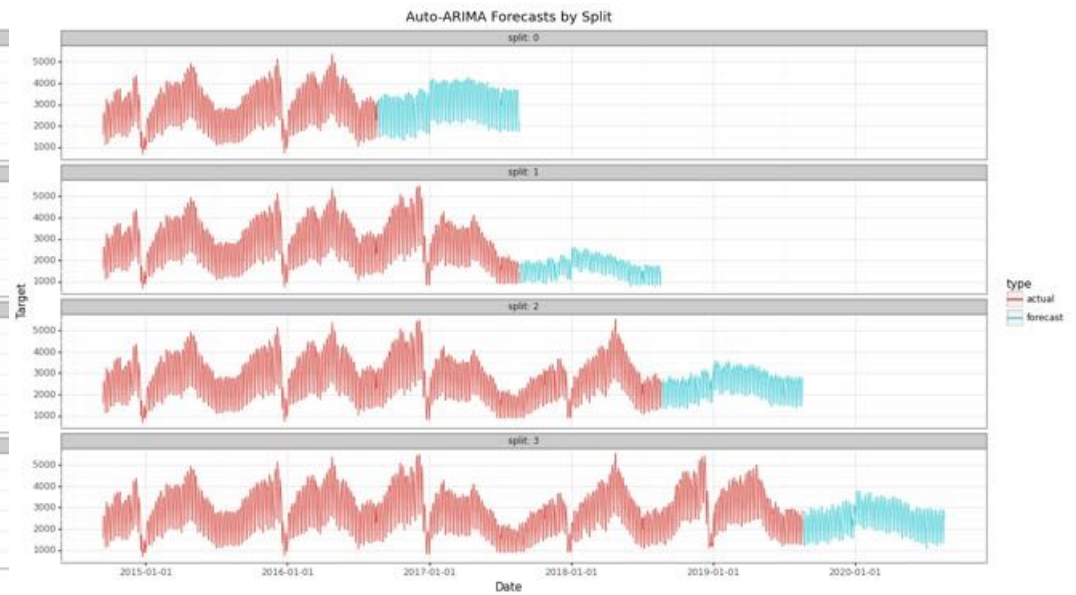

Trend Fit and Forecast by Split (trend=1095 and lags=21)



De-trended by Split



Conclusion



Method	Avg Error (across 4 splits)	Avg Train Time (across 4 splits)
Trend + Tree-Model	14.6%	~2sec
Auto-ARIMA (SARIMAX)	21.6%	~2min


The End

- Some other forecasting tools worth checking out:
 - Azure AutoML
 - Python SDK: [Set up AutoML for time-series forecasting - Azure Machine Learning | Microsoft Docs](#)
 - No Code: [Tutorial: Demand forecasting & AutoML - Azure Machine Learning | Microsoft Docs](#)
 - Modeltime (Business-Science)
 - R: <https://business-science.github.io/modeltime/>
 - Prophet
 - Github: <https://github.com/facebook/prophet>
 - Python: <https://pypi.org/project/prophet/>
 - R: <https://cran.r-project.org/web/packages/prophet/index.html>
 - Neural Prophet
 - <https://pytorch.org/>
 - Much faster than Prophet, uses PyTorch
 - Uses AR-Net for modeling autocorrelation
 - Python notebooks and PPT will be available on Lev's repo
-

Bonus: Azure AutoML Experiment: Daily Website Data Set

Algorithm name	Explained	Nor... ↑	Sampling	Submitted time	Duration	Hyperparameter
MaxAbsScaler, ExtremeRandomTrees	View explanation	0.09609	100.00 %	Apr 15, 2022 2:22 PM	4s	<code>bootstrap</code> <code>criterion : mse</code> <code>max_features : 0.8</code> <code>min_samples_leaf : 0.00236468227726</code>
MinMaxScaler, RandomForest		0.09786	100.00 %	Apr 15, 2022 11:42 AM	4s	<code>bootstrap</code> <code>criterion : mse</code> <code>max_features : 0.7</code> <code>min_samples_leaf : 0.00419663374756</code>
MinMaxScaler, RandomForest		0.09786	100.00 %	Apr 15, 2022 11:42 AM	4s	<code>bootstrap : true</code> <code>criterion : mse</code> <code>max_features : sqrt</code> <code>min_samples_leaf : 0.003466237</code>
MinMaxScaler, ExtremeRandomTrees		0.09980	100.00 %	Apr 15, 2022 11:42 AM	4s	<code>bootstrap : true</code> <code>criterion : mse</code> <code>max_features : 0.7</code> <code>min_samples_leaf : 0.0034662374</code>
StandardScalerWrapper, LightGBM		0.10009	100.00 %	Apr 15, 2022 11:42 AM	4s	<code>min_data_in_leaf : 20</code>
MinMaxScaler, ExtremeRandomTrees		0.10034	100.00 %	Apr 15, 2022 2:18 PM	4s	<code>bootstrap</code> <code>criterion : mse</code> <code>max_features : 0.5</code> <code>min_samples_leaf : 0.00508093718889</code>
TCNForecaster	Not supported	0.10238	100.00 %	Apr 15, 2022 3:17 PM	7m 13s	
StandardScalerWrapper, XGBoostRegressor		0.10365	100.00 %	Apr 15, 2022 11:42 AM	4s	<code>tree_method : auto</code>
TCNForecaster	Not supported	0.10670	100.00 %	Apr 15, 2022 11:42 AM	17m 24s	
TCNForecaster	Not supported	0.10712	100.00 %	Apr 15, 2022 11:42 AM	1h 2m 43s	
ProphetModel		0.10724	100.00 %	Apr 15, 2022 11:42 AM	4s	
TCNForecaster	Not supported	0.11375	100.00 %	Apr 15, 2022 11:42 AM	15m 20s	
RobustScaler, ExtremeRandomTrees		0.11486	100.00 %	Apr 15, 2022 2:26 PM	4s	<code>bootstrap : true</code> <code>criterion : mse</code> <code>max_features : 0.9</code> <code>min_samples_leaf : 0.0090172082</code>

AutoML Explanations: MaxAbsScaler, ExtremeRandomTrees



MaxAbsScaler,
ExtremeRandomTrees

Explained variance
0.75536

Mean absolute error
361.21

Mean absolute percentage error
12.425

Median absolute error
283.53

Normalized mean absolute error
0.074109

Normalized median absolute error
0.058172

Normalized root mean squared error
0.096089

Normalized root mean squared log error
0.083922

R2 score
0.72150

Root mean squared error
468.34

Root mean squared log error
0.17756

Spearman correlation
0.89632

StandardScalerWrapper,
LightGBM

Explained variance
0.70605

Mean absolute error
371.67

Mean absolute percentage error
12.407

Median absolute error
289.12

Normalized mean absolute error
0.076256

Normalized median absolute error
0.059318

Normalized root mean squared error
0.10009

Normalized root mean squared log error
0.083404

R2 score
0.69794

Root mean squared error
487.82

Root mean squared log error
0.17647

Spearman correlation
0.87858

ProphetModel

Explained variance
0.74682

Mean absolute error
409.74

Mean absolute percentage error
14.893

Median absolute error
327.71

Normalized mean absolute error
0.084067

Normalized median absolute error
0.067237

Normalized root mean squared error
0.10724

Normalized root mean squared log error
0.11275

R2 score
0.65320

Root mean squared error
522.71

Root mean squared log error
0.23856

Spearman correlation
0.89201

Bonus: Comparing Different Ensemble Methods

```
# from sklearn.ensemble:
    GradientBoostingRegressor
    HistGradientBoostingRegressor # (best, inspired by LightGBM)
    RandomForestRegressor
    ExtraTreesRegressor

# from xgboost !pip install xgboost
    XGBRegressor
    XGBRFRegressor
```

All regressors had similar speed and accuracy.

But the best was **HistGradientBoostingRegressor** - the histogram-based gradient boosting.

- <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.HistGradientBoostingRegressor.html>
- <https://machinelearningmastery.com/histogram-based-gradient-boosting-ensembles/>
- https://inria.github.io/scikit-learn-mooc/python_scripts/ensemble_hist_gradient_boosting.html

Why the "histogram-based" approach is faster and better?

- The bottleneck in training any boosted trees algorithm is the time for finding the best split between values when building the trees
 - This time can be significantly reduced if we pre-process the data for a given tree – aggregate them into "bins"
 - So, we convert the continuous float numbers into limited number of discrete integer values
 - This "histogram binning" reduces the number of tests needed to find best split
 - This makes finding the split much faster - but reduces the accuracy
 - To increase accuracy, we can simply increase the number of trees
-